

# Memory Locations For Variables

# A Binding Question

- Variables are bound (dynamically) to values
- Those values must be stored somewhere
- Therefore, variables must somehow be bound to memory locations
- How?

# Functional Meets Imperative

- Imperative languages expose the concept of memory locations:  **$a := 0$** 
  - Store a zero in  **$a$** 's memory location
- Functional languages hide it:  **$\text{val } a = 0$** 
  - Bind  **$a$**  to the value zero
- But both need to connect variables to values represented in memory
- So both face the same binding question

# Outline

- Activation records
- Static allocation of activation records
- Stacks of activation records
- Handling nested function definitions
- Functions as parameters
- Long-lived activation records

# Function Activations

- The lifetime of one execution of a function, from call to corresponding return, is called an *activation* of the function
- When each activation has its own binding of a variable to a memory locations, it is an *activation-specific* variable
- (Also called *dynamic* or *automatic*)

# Activation-Specific Variables

- In most modern languages, activation-specific variables are the most common kind:

```
fun days2ms days =  
  let  
    val hours = days * 24.0  
    val minutes = hours * 60.0  
    val seconds = minutes * 60.0  
  in  
    seconds * 1000.0  
  end;
```

# Block Activations

- For block constructs that contain code, we can speak of an activation of the *block*
- The lifetime of one execution of the block
- A variable might be specific to an activation of a particular block within a function:

```
fun fact n =  
  if (n=0) then 1  
  else let val b = fact (n-1) in n*b  
end;
```

# Other Lifetimes For Variables

- Most imperative languages have a way to declare a variable that is bound to a single memory location for the entire runtime
- Obvious binding solution: static allocation (classically, the loader allocates these)

```
int count = 0;
int nextcount() {
    count = count + 1;
    return count;
}
```



# Scope And Lifetime Differ

- In most modern languages, variables with local *scope* have activation-specific *lifetimes*, at least by default
- However, these two aspects can be separated, as in C:

```
int nextcount() {  
    static int count = 0;  
    count = count + 1;  
    return count;  
}
```

# Other Lifetimes For Variables

- Object-oriented languages use variables whose lifetimes are associated with object lifetimes
- Some languages have variables whose values are persistent: they last across multiple executions of the program
- Today, we will focus on activation-specific variables

# Activation Records

- Language implementations usually allocate all the activation-specific variables of a function together as an *activation record*
- The activation record also contains other activation-specific data, such as
  - Return address: where to go in the program when this activation returns
  - Link to caller's activation record: more about this in a moment

# Block Activation Records

- When a block is entered, space must be found for the local variables of that block
- Various possibilities:
  - Preallocate in the containing function's activation record
  - Extend the function's activation record when the block is entered (and revert when exited)
  - Allocate separate block activation records
- Our illustrations will show the first option

# Outline

- Activation-specific variables
- **Static allocation of activation records**
- Stacks of activation records
- Handling nested function definitions
- Functions as parameters
- Long-lived activation records

# Static Allocation

- The simplest approach: allocate one activation record for every function, statically
- Older dialects of Fortran and Cobol used this system
- Simple and fast

# Example

```
FUNCTION AVG (ARR, N)
DIMENSION ARR(N)
SUM = 0.0
DO 100 I = 1, N
    SUM = SUM + ARR(I)
100 CONTINUE
AVG = SUM / FLOAT(N)
RETURN
END
```

N address
ARR address
return address
I
SUM
AVG

# Drawback

- Each function has one activation record
- There can be only one activation alive at a time
- Modern languages (including modern dialects of Cobol and Fortran) do not obey this restriction:
  - Recursion
  - Multithreading



# Outline

- Activation-specific variables
- Static allocation of activation records
- **Stacks of activation records**
- Handling nested function definitions
- Functions as parameters
- Long-lived activation records

# Stacks Of Activation Records

- To support recursion, we need to allocate a new activation record for each activation
- Dynamic allocation: activation record allocated when function is called
- For many languages, like C, it can be deallocated when the function returns
- A stack of activation records: *stack frames* pushed on call, popped on return


# Current Activation Record

- Before, static: location of activation record was determined before runtime
- Now, dynamic: location of the *current* activation record is not known until runtime
- A function must know how to find the address of its current activation record
- Often, a machine register is reserved to hold this

# C Example

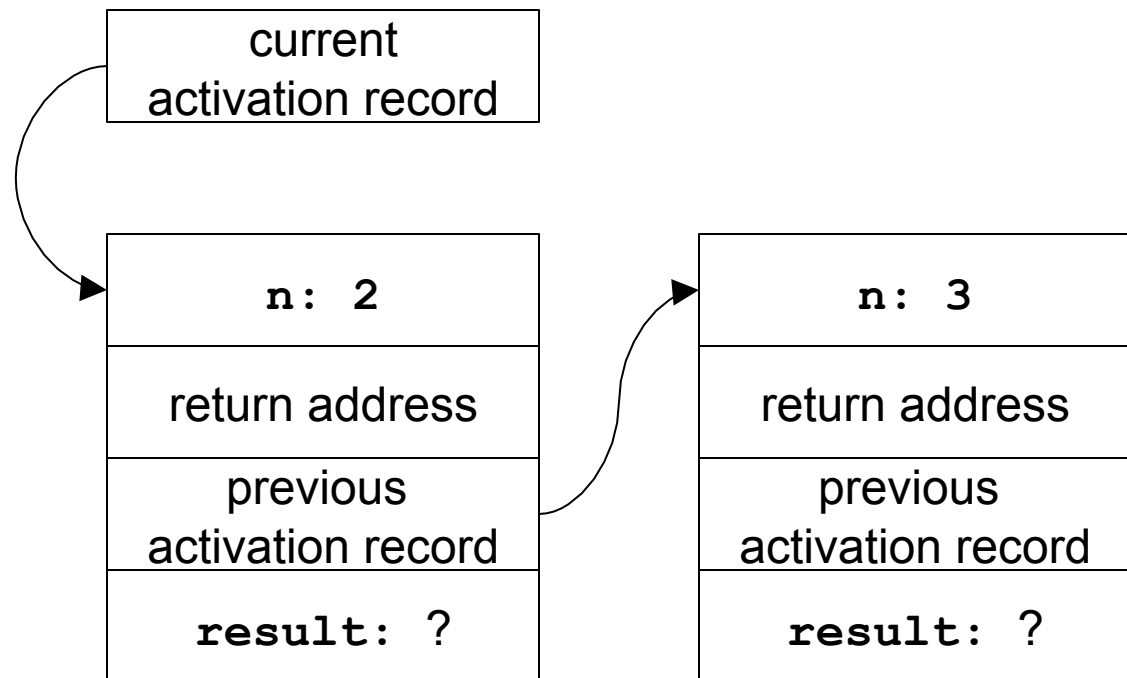
*We are evaluating **fact(3)**. This shows the contents of memory just before the recursive call that creates a second activation.*

```
int fact(int n) {  
    int result;  
    if (n<2) result = 1;  
    else result = n * fact(n-1);  
    return result;  
}
```



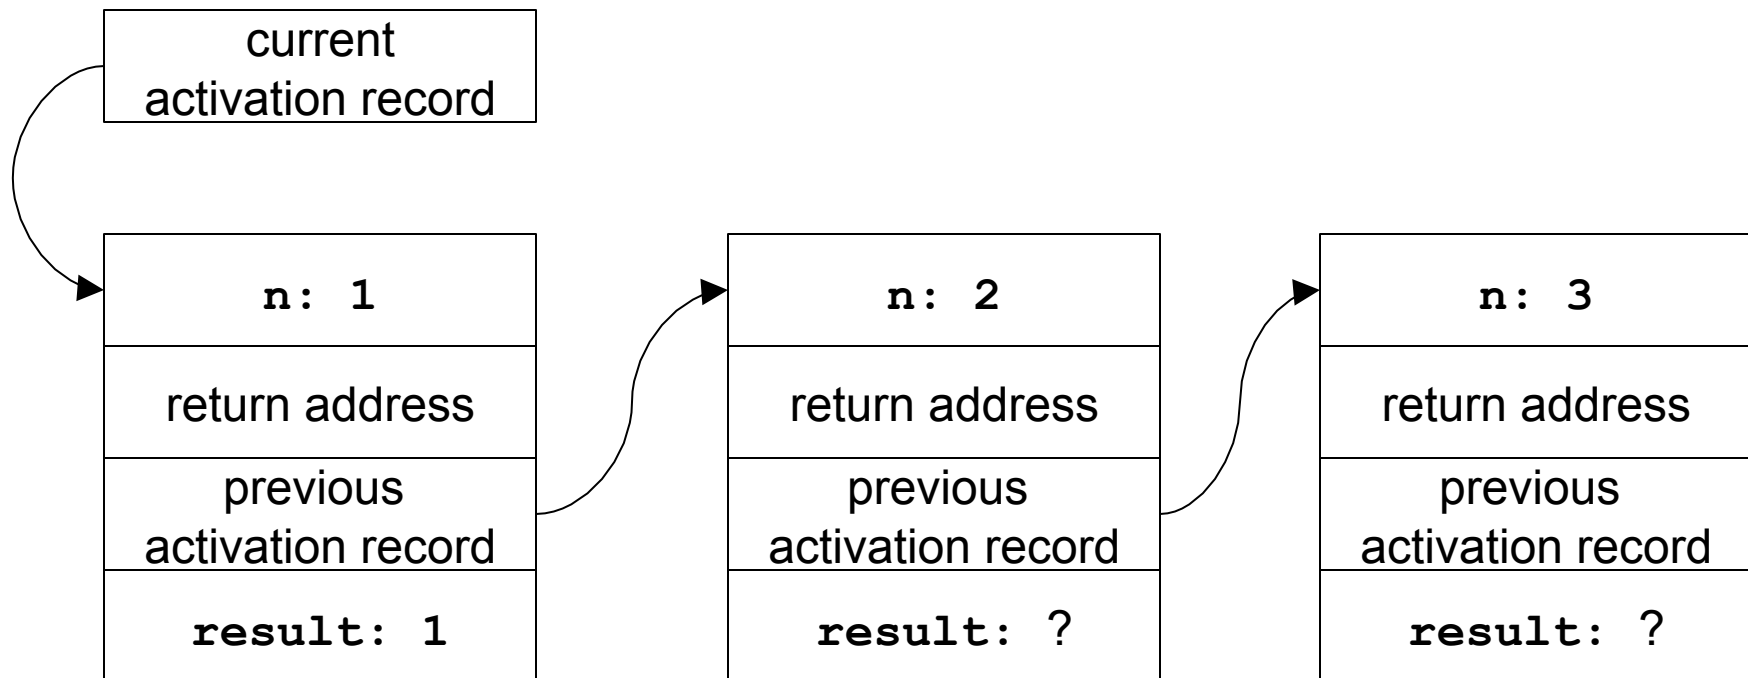
*This shows the contents of memory just before the third activation.*

```
int fact(int n) {  
    int result;  
    if (n<2) result = 1;  
    else result = n * fact(n-1);  
    return result;  
}
```



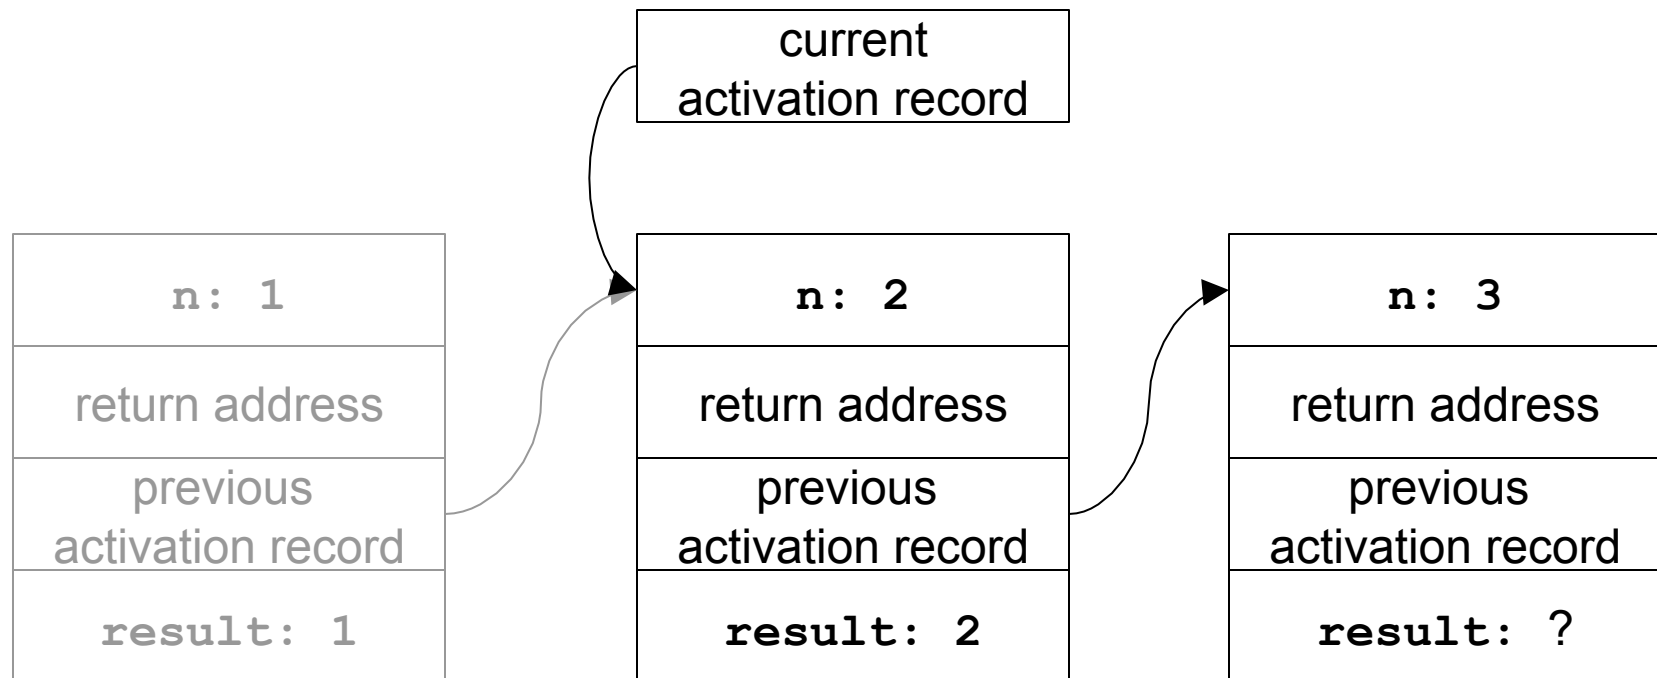
*This shows the contents of memory just before the third activation returns.*

```
int fact(int n) {  
    int result;  
    if (n<2) result = 1;  
    else result = n * fact(n-1);  
    return result;  
}
```



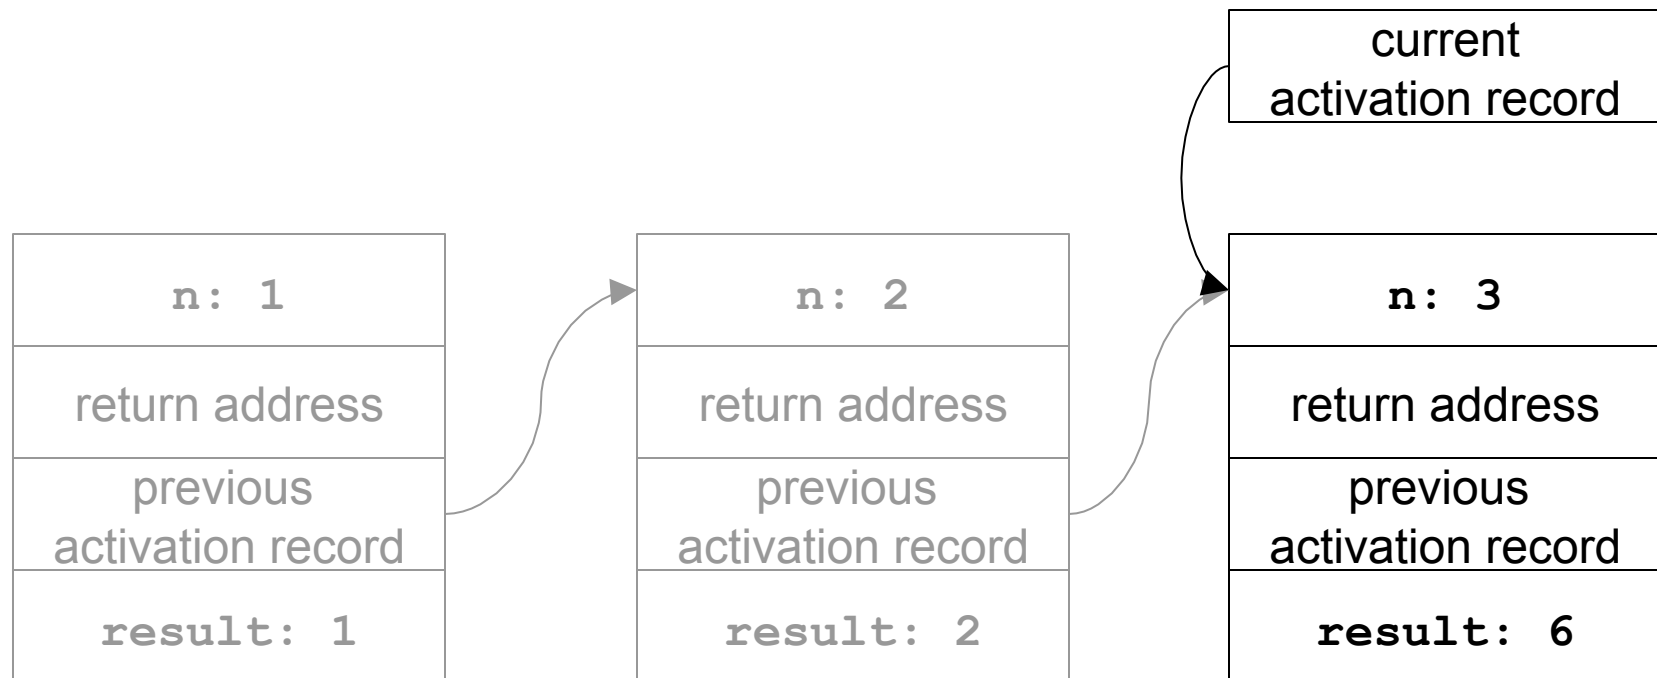
*The second activation is about to return.*

```
int fact(int n) {  
    int result;  
    if (n<2) result = 1;  
    else result = n * fact(n-1);  
    return result;  
}
```



*The first activation is about to return with the result* **fact(3) = 6**.

```
int fact(int n) {  
    int result;  
    if (n<2) result = 1;  
    else result = n * fact(n-1);  
    return result;  
}
```





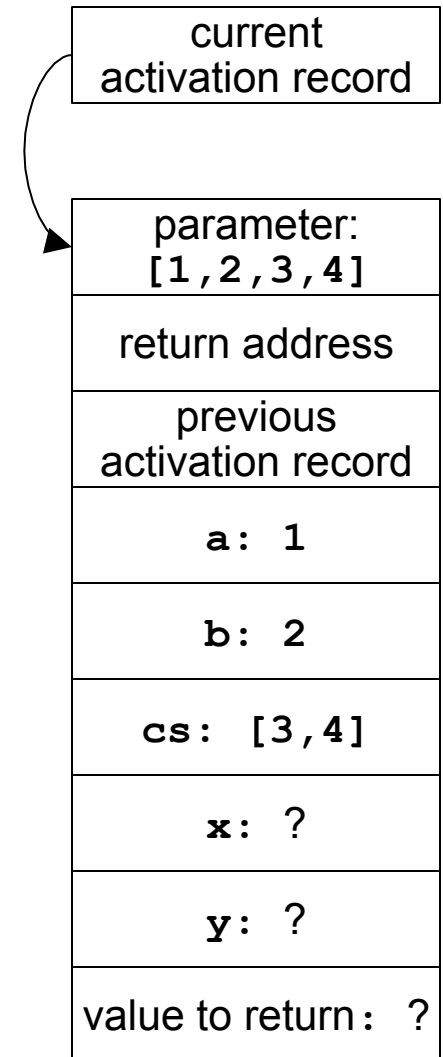
# ML Example

*We are evaluating*

**halve** [1,2,3,4].

*This shows the contents of  
memory just before the  
recursive call that creates  
a second activation.*

```
fun halve nil = (nil, nil)
|   halve [a] = ([a], nil)
|   halve (a::b::cs) =
    let
      val (x, y) = halve cs
    in
      (a::x, b::y)
    end;
```

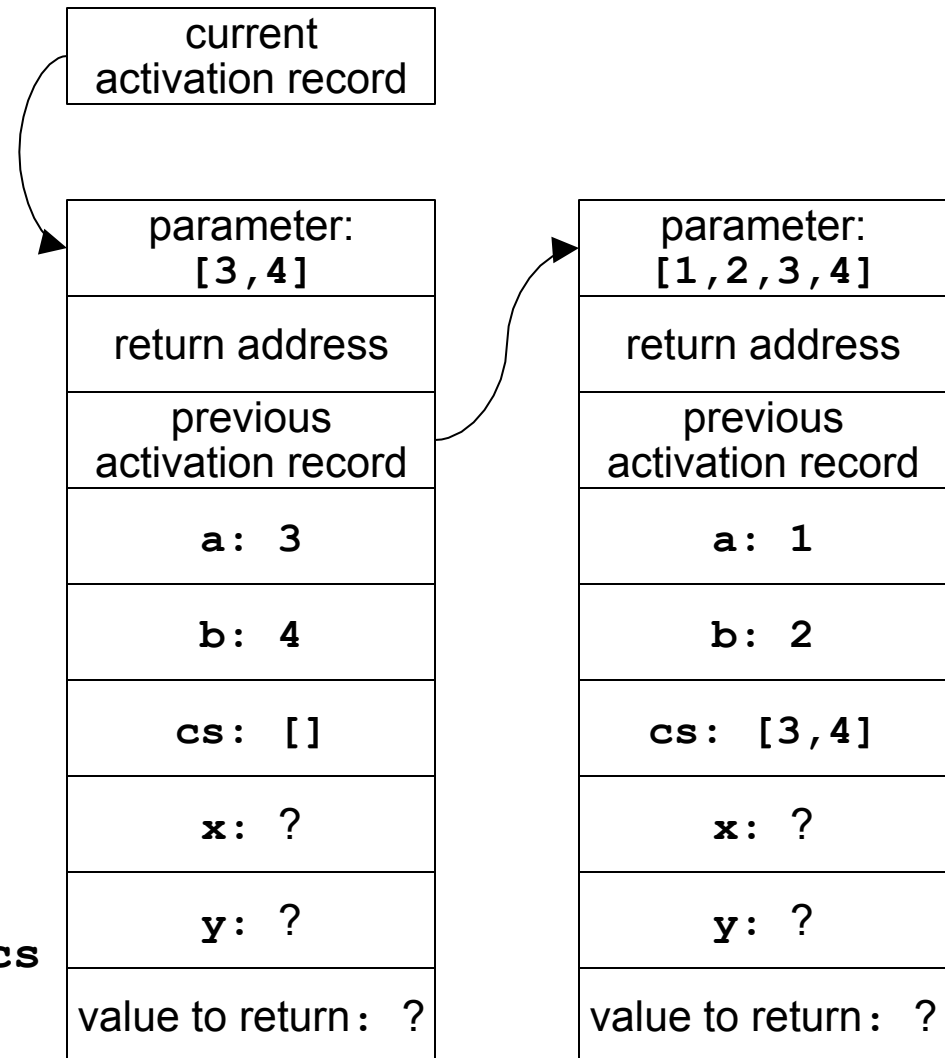


*This shows the contents of memory just before the third activation.*

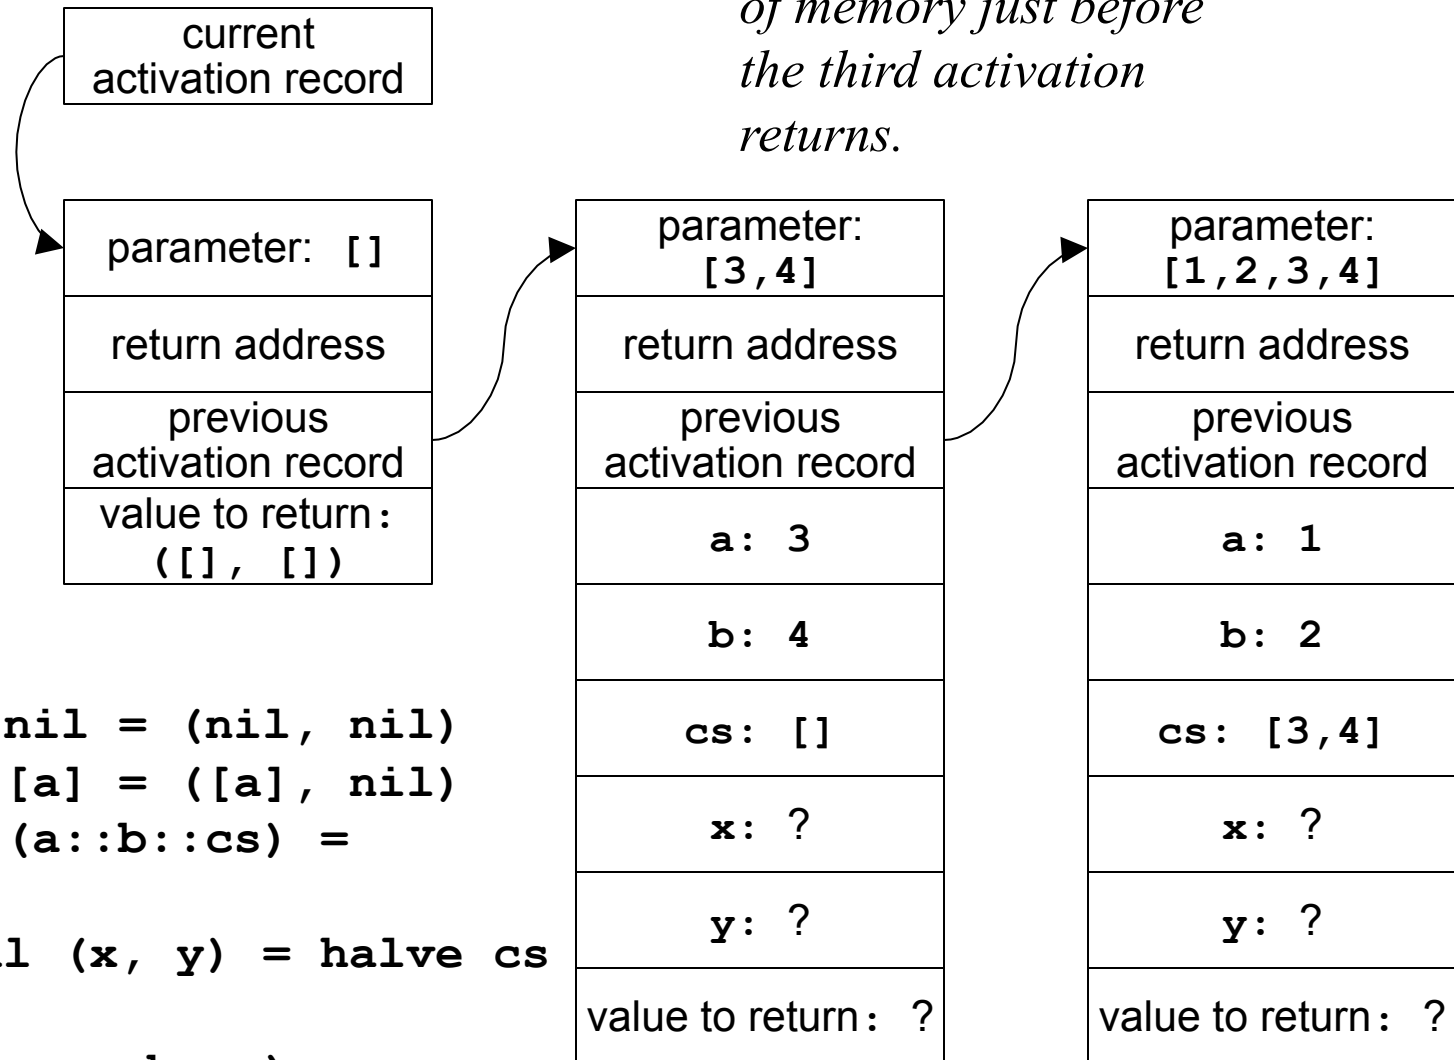
```

fun halve nil = (nil, nil)
|   halve [a] = ([a], nil)
|   halve (a::b::cs) =
    let
      val (x, y) = halve cs
    in
      (a::x, b::y)
    end;

```



*This shows the contents  
of memory just before  
the third activation  
returns.*

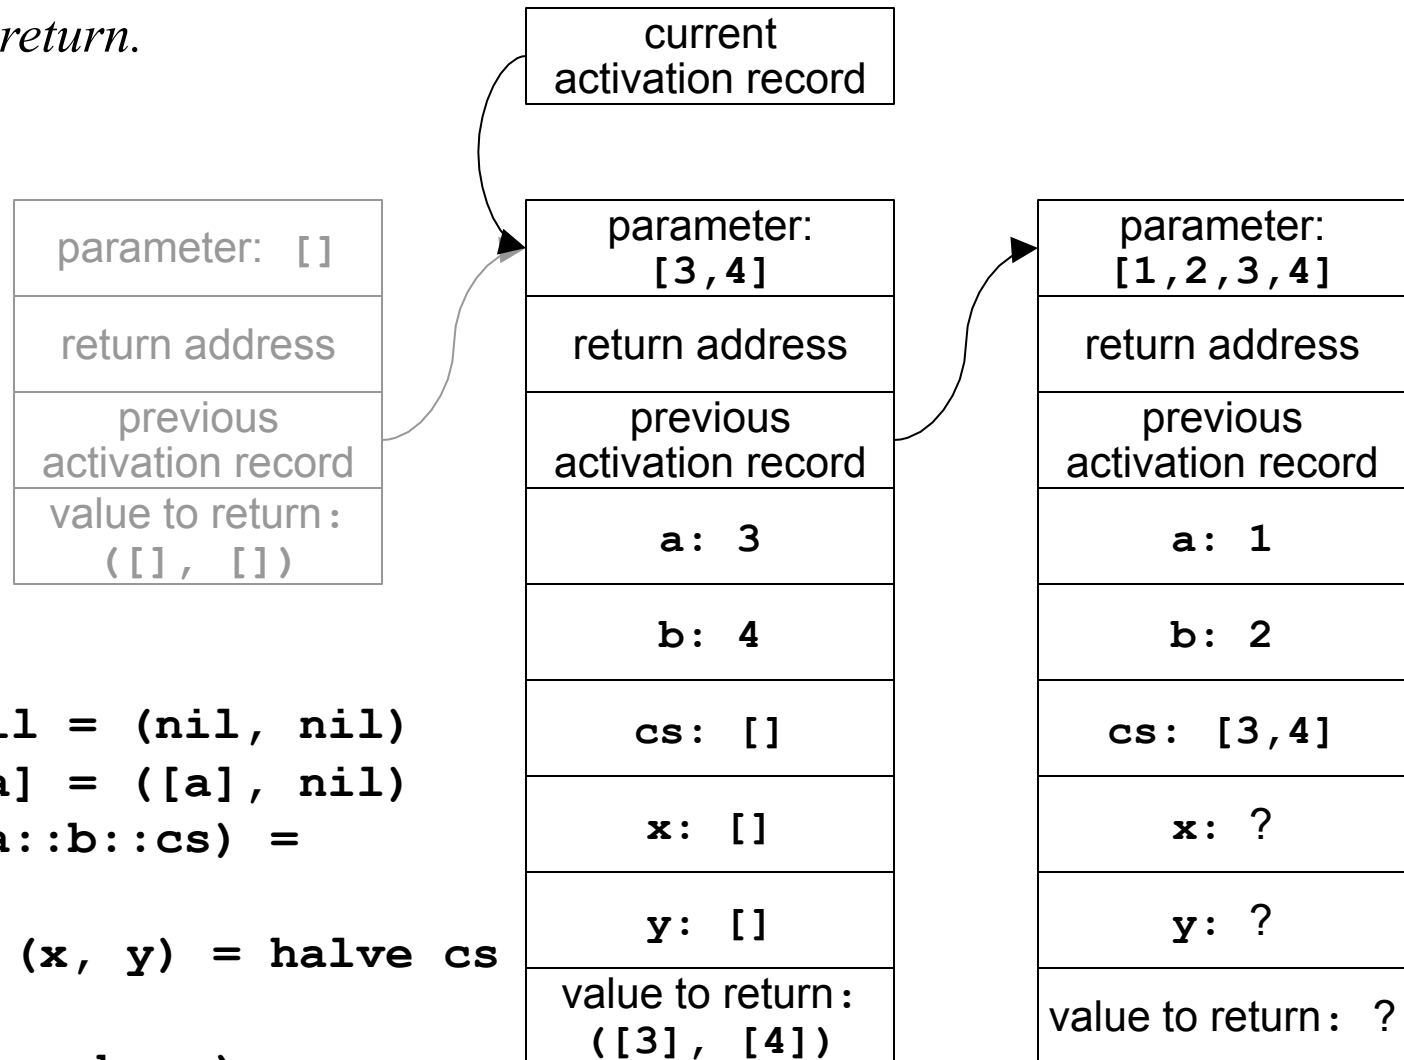


```

fun halve nil = (nil, nil)
|   halve [a] = ([a], nil)
|   halve (a::b::cs) =
    let
      val (x, y) = halve cs
    in
      (a::x, b::y)
    end;

```

*The second activation is about to return.*



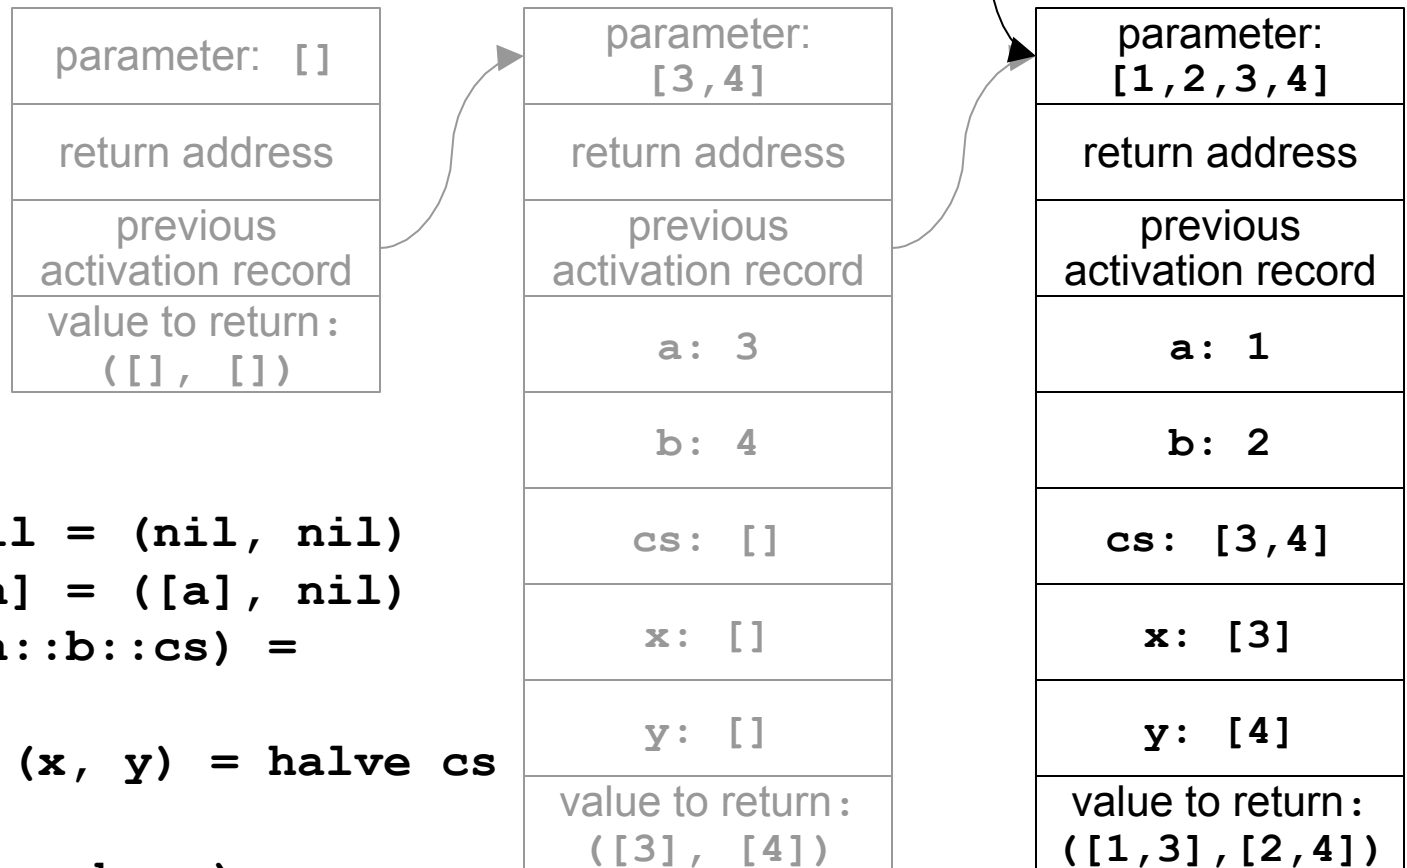
```

fun halve nil = (nil, nil)
|   halve [a] = ([a], nil)
|   halve (a::b::cs) =
    let
      val (x, y) = halve cs
    in
      (a::x, b::y)
    end;

```

*The first activation is about to return with the result*

**halve** [1,2,3,4] =  
([1,3], [2,4])



```
fun halve nil = (nil, nil)
|   halve [a] = ([a], nil)
|   halve (a::b::cs) =
    let
      val (x, y) = halve cs
    in
      (a::x, b::y)
    end;
```

# Outline


- Activation-specific variables
- Static allocation of activation records
- Stacks of activation records
- **Handling nested function definitions**
- Functions as parameters
- Long-lived activation records

# Nesting Functions

- What we just saw is adequate for many languages, including C
- But not for languages that allow this trick:
  - Function definitions can be nested inside other function definitions
  - Inner functions can refer to local variables of the outer functions (under the usual block scoping rule)
- Like ML, Ada, Pascal, etc.

# Example

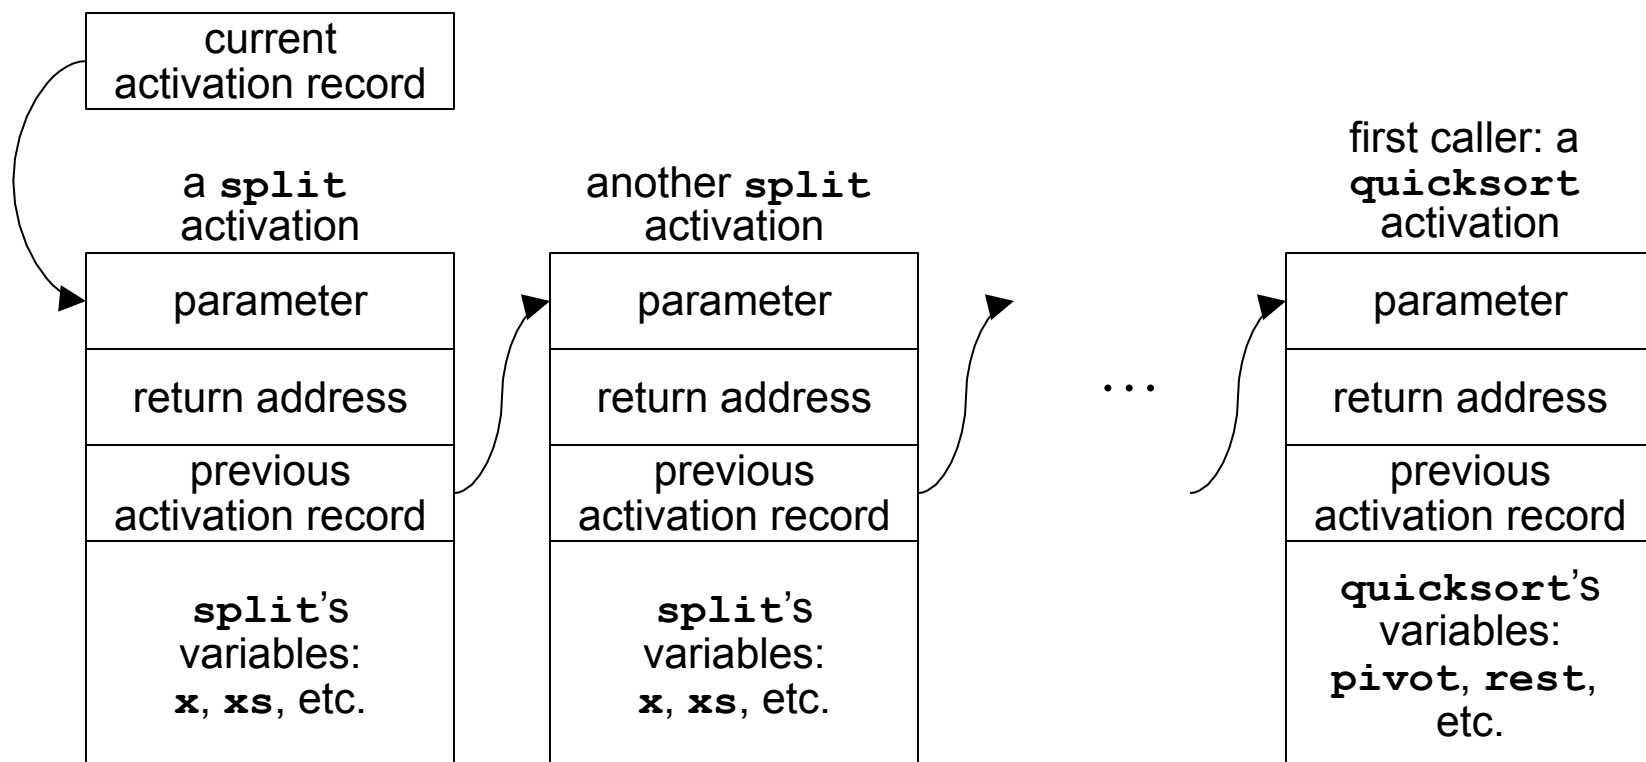
```
fun quicksort nil = nil
| quicksort (pivot::rest) =
  let
    fun split(nil) = (nil,nil)
    | split(x::xs) =
      let
        val (below, above) = split(xs)
      in
        if x < pivot then (x::below, above)
        else (below, x::above)
      end;
    val (below, above) = split(rest)
  in
    quicksort below @ [pivot] @ quicksort above
  end;
```






# The Problem

- How can an activation of the inner function (**split**) find the activation record of the outer function (**quicksort**)?
- It isn't necessarily the previous activation record, since the caller of the inner function may be another inner function
- Or it may call itself recursively, as **split** does...



# Nesting Link


- An inner function needs to be able to find the address of the most recent activation for the outer function
- We can keep this *nesting link* in the activation record...



# Setting The Nesting Link

- Easy if there is only one level of nesting:
  - Calling outer function: set to null
  - Calling from outer to inner: set nesting link same as caller's activation record
  - Calling from inner to inner: set nesting link same as caller's nesting link
- More complicated if there are multiple levels of nesting...

# Multiple Levels Of Nesting



- References at the same level (**f1** to **v1**, **f2** to **v2**, **f3** to **v3** ) use current activation record
- References  $n$  nesting levels away chain back through  $n$  nesting links

# Other Solutions

- The problem: references from inner functions to variables in outer ones
  - Nesting links in activation records: as shown
  - Displays: nesting links not in the activation records, but collected in a single static array
  - Lambda lifting: problem references replaced by references to new, hidden parameters

# Outline

- Activation-specific variables
- Static allocation of activation records
- Stacks of activation records
- Handling nested function definitions
- **Functions as parameters**
- Long-lived activation records



# Functions As Parameters

- When you pass a function as a parameter, what really gets passed?
- Code must be part of it: source code, compiled code, pointer to code, or implementation in some other form
- For some languages, something more is required...

# Example

```
fun addXToAll (x,theList) =  
  let  
    fun addX y =  
      y + x;  
  in  
    map addX theList  
  end;
```


- This function adds **x** to each element of **theList**
- Notice: **addXToAll** calls **map**, **map** calls **addX**, and **addX** refers to a variable **x** in **addXToAll**'s activation record

# Nesting Links Again

- When **map** calls **addX**, what nesting link will **addX** be given?
  - Not **map**'s activation record: **addX** is not nested inside **map**
  - Not **map**'s nesting link: **map** is not nested inside anything
- To make this work, the parameter **addX** passed to **map** must include the nesting link to use when **addX** is called


# Not Just For Parameters

- Many languages allow functions to be passed as parameters
- Functional languages allow many more kinds of operations on function-values:
  - passed as parameters, returned from functions, constructed by expressions, etc.
- Function-values include both parts: code to call, and nesting link to use when calling it



# Example

```
fun addXToAll (x,theList) =  
  let  
    fun addX y =  
      y + x;  
  in  
    map addX theList  
  end;
```



*This shows the contents of memory just before the call to **map**. The variable **addX** is bound to a function-value including code and nesting link.*

# Outline

- Activation-specific variables
- Static allocation of activation records
- Stacks of activation records
- Handling nested function definitions
- Functions as parameters
- Long-lived activation records

# One More Complication

- What happens if a function value is used after the function that created it has returned?

```
fun test () =  
  let  
    val f = funToAddX 3;  
  in  
    f 5  
  end;
```

```
fun funToAddX x =  
  let  
    fun addX y =  
      y + x;  
  in  
    addX  
  end;
```

*Note: **test**'s parameter here is the special value (). That's the one and only value of type **unit** in ML. It often serves as a dummy parameter—a sort of placeholder for functions that don't have significant parameters.*

```
fun test () =
```

```
  let
```

```
    val f = funToAddX 3;
```

```
  in
```

```
    f 5
```

```
  end;
```

```
fun funToAddX x =
```

```
  let
```

```
    fun addX y =
```

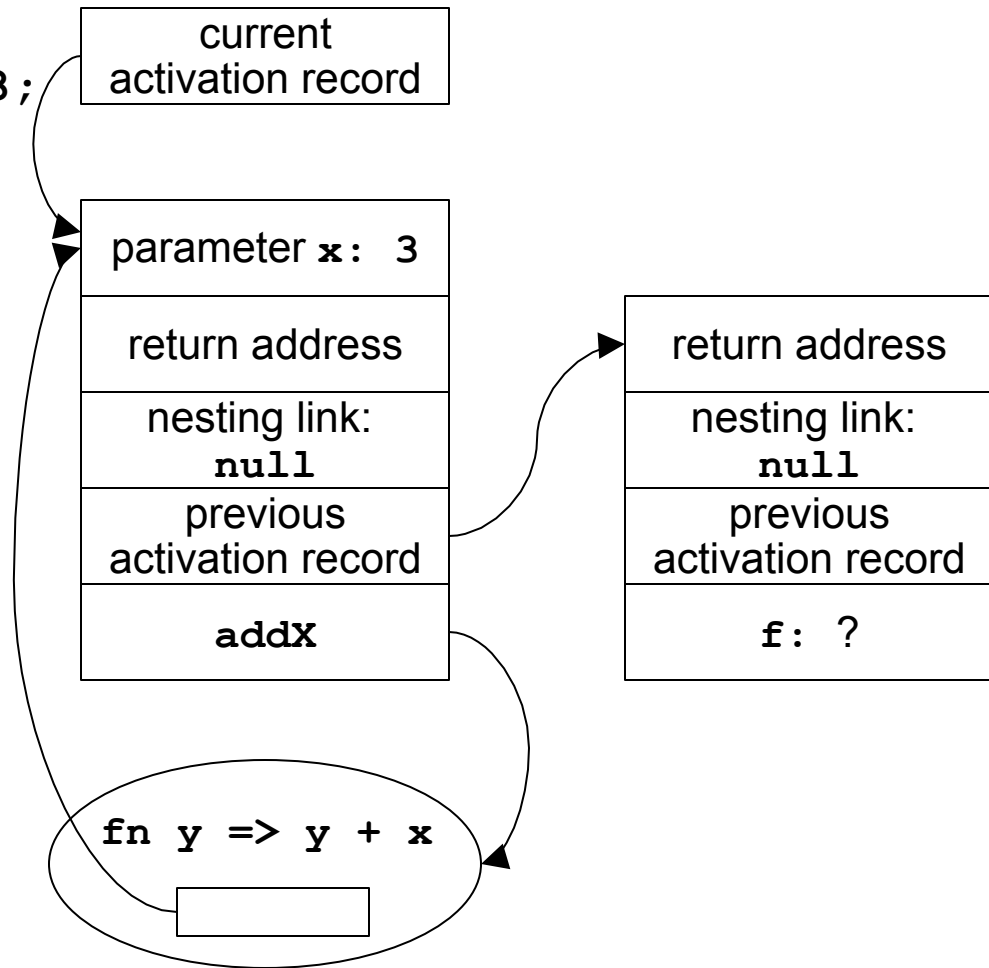
```
      y + x;
```

```
  in
```

```
    addX
```

```
  end;
```

*This shows the contents of  
memory just before  
**funToAddX** returns.*





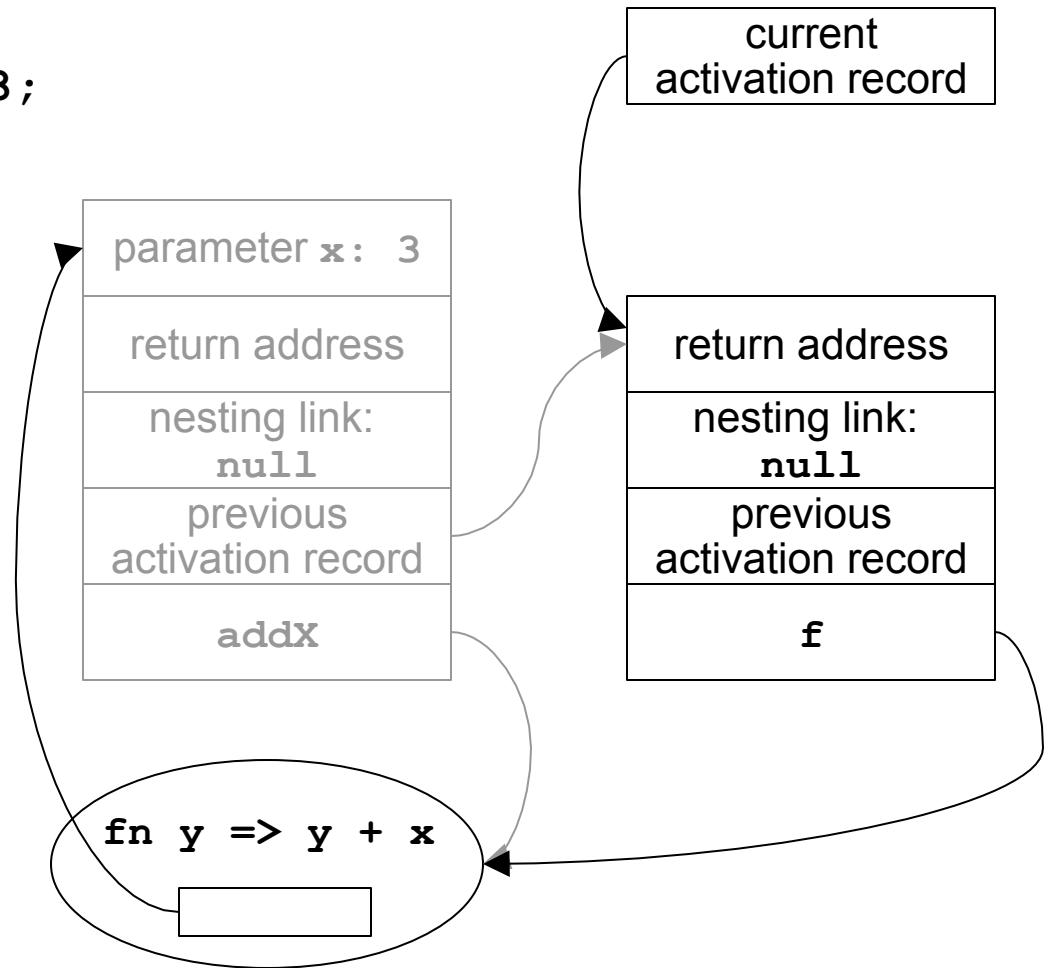
```

fun test () =
  let
    val f = funToAddX 3;
  in
    f 5
  end;

fun funToAddX x =
  let
    fun addX y =
      y + x;
  in
    addX
  end;

```

*After **funToAddX** returns, **f** is the bound to the new function-value.*



# The Problem

- When **test** calls **f**, the function will use its nesting link to access **x**
- That is a link to an activation record for an activation that is finished
- This will fail if the language system deallocated that activation record when the function returned

# The Solution

- For ML, and other languages that have this problem, activation records cannot always be allocated and deallocated in stack order
- Even when a function returns, there may be links to its activation record that will be used; it can't be deallocated it is unreachable
- *Garbage collection*: chapter 14, coming soon!

# Conclusion

- The more sophisticated the language, the harder it is to bind activation-specific variables to memory locations
  - Static allocation: works for languages that permit only one activation at a time (like early dialects of Fortran and Cobol)
  - Simple stack allocation: works for languages that do not allow nested functions (like C)

# Conclusion, Continued

- Nesting links (or some such trick): required for languages that allow nested functions (like ML, Ada and Pascal); function values must include both code and nesting link
- Some languages (like ML) permit references to activation records for activations that are finished; so activation records cannot be deallocated on return